

Ali Hamed

AI Engineer | LLMs, Agents, RAG, PyTorch, Python, Production AI

ali.h.10j@gmail.com • +972-53-339-2826 • Haifa District, Israel • [GitHub](#) • [LinkedIn](#) • [Portfolio](#)
Languages: Arabic & Hebrew (native) • English (professional)

SUMMARY

AI Engineer with **2 years** of hands-on industry experience taking AI from prototype to production — building and shipping LLM agents, RAG pipelines, and ML models that run in live engineering workflows and deliver measurable impact. Owns the full lifecycle: data pipelines, model selection and fine-tuning, retrieval and tool-calling architecture, evaluation, API serving, containerized deployment, and monitoring. Combines applied AI with strong software and infrastructure engineering (Python, REST/FastAPI, Docker, Kubernetes, GPU/CUDA, CI/CD) on complex distributed 4G/5G Open RAN systems. M.Sc. Information Systems & AI (GPA 93). **NVIDIA AI Infrastructure** and **IBM Generative AI & LLMs** certified.

EXPERIENCE

Parallel Wireless — AI Automation Engineer	09.2024 – present · 2 years
- Delivered 3 production AI agents end to end (Python, LangChain, MCP tool calling, FAISS). - Own the full lifecycle: data pipeline, retrieval, evaluation, serving, deployment, monitoring. - Designed RAG over telemetry: chunking, embeddings, FAISS indexing, and semantic retrieval. - Implemented tool calling so agents fetch context and trigger actions, with human-in-the-loop. - Fine-tuned HuggingFace transformers for log summarization and failure classification. - Built anomaly-detection and classification models (Scikit-learn) evaluated with ROC/AUC/F1. - Engineered ETL pipelines (Pandas, NumPy, SQL, Bash) for training- and inference-ready datasets. - Served models over REST/WebSocket APIs (FastAPI) on Docker/Kubernetes with GPU/CUDA. - Tracked experiments and versioned artifacts with Git, MLflow, and Weights & Biases. - Built evaluation datasets and regression suites to prevent agent-quality regressions. - Cut manual triage: agents summarize CI logs, classify failures, and route incidents. - Grounded AI in production engineering: Jenkins CI/CD, 200–500 tests, 8h → 2h (~75%). - Accelerated development with Claude Code and Cursor; collaborated cross-functionally in Agile.	

EDUCATION

M.Sc., Information Systems & AI • University of Haifa	GPA 93 • Expected 2027
Specialization in Artificial Intelligence, Machine Learning, Deep Learning, NLP, Big Data Systems, and Software Engineering; research-oriented coursework with applied projects.	
B.Sc., Computer Science • Ruppin Academic Center	GPA 83 • 2022 – 2025
Specialization in AI & Data Science. Coursework in Software Development, Data Science, Machine Learning, and Deep Learning; strong OOP, data structures, algorithms, and databases.	

CERTIFICATIONS

NVIDIA AI Infrastructure and Operations Fundamentals • Sep 2025	Certificate
IBM Generative AI and LLMs • Apr 2025	Certificate
NVIDIA Networking • Apr 2025	Certificate
AWS Cloud Developing • Dec 2024	Certificate
AWS Cloud Practitioner Essentials • Nov 2024	Certificate

SKILLS

LLMs & Agents	LLMs (OpenAI, Claude, Llama) · LangChain · LangGraph · autonomous AI agents · multi-agent systems · MCP servers / tool calling · agentic workflows · LLM orchestration · prompt engineering · context engineering · guardrails · human-in-the-loop
RAG & Retrieval	RAG pipelines · FAISS / vector databases · embeddings · semantic search · chunking strategies · hybrid retrieval · context windows
ML & Deep Learning	PyTorch · TensorFlow · HuggingFace · Scikit-learn · Transformers · CNNs · Autoencoders · NLP · Computer Vision · fine-tuning (LoRA/PEFT) · transfer learning · Random Forest · ensembles
Production AI / MLOps	model deployment · model serving · inference pipelines · model monitoring · experiment tracking (MLflow, W&B) · evaluation frameworks · error analysis · A/B evaluation · GPU/CUDA · LLMOps
Engineering & Data	Python · FastAPI · REST / WebSocket APIs · SQL · Pandas · NumPy · Bash · ETL · data pipelines · Git · CI/CD (Jenkins) · Docker · Kubernetes · AWS (SageMaker, EC2, S3, Lambda) · Linux
Practices & Methodologies	end-to-end AI delivery (prototype → production) · ML lifecycle · MLOps · LLMOps · data preprocessing · model evaluation & validation · Agile · code review · Git workflows · technical documentation · stakeholder communication

VOLUNTEER

Volunteer Instructor • Hezkek LaAtid "Scientific Excellence", Fureidis High School	09.2025 – present
Teaching Computer Science fundamentals and programming to 40 students.	
Student Tutor • Ruppin Academic Center (Algebra, Python, Discrete Math) "Perach" Program • Tutor & mentor for elementary school students	01.2023 – 02.2024 2022 – 2025